

Lab1: RV64内核引导与时钟中断处理

张志心 3210106357

1 实验目的

- 学习 RISC-V 汇编，编写 `head.S` 实现跳转到内核运行的第一个 C 函数。
- 学习 OpenSBI，理解 OpenSBI 在实验中所起到的作用，并调用 OpenSBI 提供的接口完成字符的输出。
- 学习 Makefile 相关知识，补充项目中的 Makefile 文件，来完成对整个工程的管理。
- 学习 RISC-V 的 `trap` 处理相关寄存器与指令，完成对 `trap` 处理的初始化。
- 理解 CPU 上下文切换机制，并正确实现上下文切换功能。
- 编写 `trap` 处理函数，完成对特定 `trap` 的处理。
- 调用 OpenSBI 提供的接口，完成对时钟中断事件的设置。

2 实验环境

- 本次实验环境为 Mac 系统下 Ubuntu22.04 VM。（不同于 Lab0）。
- 环境配置过程与 Lab0 类似。

3 实验步骤

3.1 RV64 内核引导

3.1.1 编写 head.S

为即将运行的第一个 C 函数设置程序栈，大小为 4kb，将栈放置在 `.bss.stack` 段。然后通过跳转指令，跳转至 `main.c` 中的 `start_kernel`。

```
# head.S
.extern start_kernel

.section .text.entry
.globl _start
_start:
# -----
# - your code here -
# -----
la sp, boot_stack_top
jal start_kernel

.section .bss.stack
.globl boot_stack
boot_stack:
.space 4096 # <-- change to your stack size
```

```
.globl boot_stack_top
boot_stack_top:
```

3.1.2 完善 Makefile 脚本

补充 lib/Makefile 如下:

```
SRCS = $(shell find . -name "*.S") $(shell find . -name "*.c")
OBJS = $(addsuffix .o, $(basename $(SRCS)))

all: $(OBJS)

%.o: %.S
    @echo CC $< $@
    @$$(GCC) $(CFLAG) -c $< -o $@

%.o: %.c
    @echo CC $< $@
    @$$(GCC) $(CFLAG) -c $< -o $@

clean:
    -@rm *.o 2>/dev/null
```

3.1.3 补充 sbi.c

1. 将 ext (Extension ID) 放入寄存器 a7 中, fid (Function ID) 放入寄存器 a6 中, 将 arg0 ~ arg5 放入寄存器 a0 ~ a5 中。
2. 使用 ecall 指令。ecall 之后系统会进入 M 模式, 之后 OpenSBI 会完成相关操作。
3. OpenSBI 的返回结果会存放在寄存器 a0, a1 中, 其中 a0 为 error code, a1 为返回值, 我们用 sbiret 来接受这两个返回值。

```
struct sbiret result;
__asm__ volatile (
    "\tmv a7, %[ext]\n"
    "\tmv a6, %[fid]\n"
    "\tmv a0, %[arg0]\n"
    "\tmv a1, %[arg1]\n"
    "\tmv a2, %[arg2]\n"
    "\tmv a3, %[arg3]\n"
    "\tmv a4, %[arg4]\n"
    "\tmv a5, %[arg5]\n"
    "\tecall\n"
    "\tmv %[err], a0\n"
    "\tmv %[res], a1\n"
    : [err] "=r" (result.error), [res] "=r" (result.value)
    : [ext] "r" (ext), [fid] "r" (fid),
      [arg0] "r" (arg0), [arg1] "r" (arg1), [arg2] "r" (arg2), [arg3] "r"
(arg3), [arg4] "r" (arg4), [arg5] "r" (arg5)
    );
return result;
```

3.1.4 修改 defs

```
#define csr_read(csr) \
({ \
    uint64 __v; \
    asm volatile("csrr %[v], " #csr \
                : [v] "=r" (__v) : : "memory"); \
    __v; \
})
```

3.1.5 qemu 运行 make 得到内核

在 /lab1 目录下进行 make

3.2 RV64 时钟中断处理

- 准备工作

```
# vmlinux.lds
.text : ALIGN(0x1000){
    _stext = .;

    *(.text.init)      <- 加入了 .text.init
    *(.text.entry)     <- 之后我们实现 中断处理逻辑 会放置在 .text.entry
    *(.text .text.*)

    _etext = .;
}
# head.S
.extern start_kernel

# .section .text.entry
.section .text.init      <- 将 _start 放入 .text.init section
.globl _start
```

3.2.1 开启 trap 处理

1. 设置 `stvec`，将 `_traps` 所表示的地址写入 `stvec`，这里我们采用 `Direct` 模式，而 `_traps` 则是 `trap` 处理入口函数的基地址。
2. 开启时钟中断，将 `sie[STIE]` 置 1。
3. 设置第一次时钟中断，参考 `clock_set_next_event()` (`clock_set_next_event()` 在 4.3.4 中介绍) 中的逻辑用汇编实现。
4. 开启 S 态下的中断响应，将 `sstatus[SIE]` 置 1。

```
# set stvec = _traps
la t0, _traps
csrw stvec, t0

# enable supervisor time interrupt
csrr t0, sie
ori t1, t0, 0x00000020
csrw sie, t1

# set first time interrupt
rdtime t0
li a0, 10000000
add a0, t0, a0
li a1, 0
li a2, 0
li a3, 0
li a4, 0
li a5, 0
li a6, 0
li a7, 0
ecall
```

```
# set sstatus[SIE] = 1
csrr t0, sstatus
ori t1, t0, 0x00000002
csrw sstatus, t1
```

3.2.2 实现上下文切换

使用汇编实现上下文切换机制，包含以下几个步骤：

1. 在 `arch/riscv/kernel/` 目录下添加 `entry.S` 文件。
2. 保存 CPU 的寄存器（上下文）到内存中（栈上）。
3. 将 `scause` 和 `sepc` 中的值传入 `trap` 处理函数 `trap_handler`（`trap_handler` 在 4.3.3 中介绍），我们将会在 `trap_handler` 中实现对 `trap` 的处理。
4. 在完成对 `trap` 的处理之后，我们从内存中（栈上）恢复 CPU 的寄存器（上下文）。
5. 从 `trap` 中返回。

```
.section .text.entry

#define context_size 256
.align 2
.globl _traps
_traps:

# allocate context struct
addi sp, sp, -context_size

sd x1, 8(sp)
sd x3, 24(sp)
sd x4, 32(sp)
sd x5, 40(sp)
sd x6, 48(sp)
sd x7, 56(sp)
sd x8, 64(sp)
sd x9, 72(sp)
sd x10, 80(sp)
sd x11, 88(sp)
sd x12, 96(sp)
sd x13, 104(sp)
sd x14, 112(sp)
sd x15, 120(sp)
sd x16, 128(sp)
sd x17, 136(sp)
sd x18, 144(sp)
sd x19, 152(sp)
sd x20, 160(sp)
sd x21, 168(sp)
sd x22, 176(sp)
sd x23, 184(sp)
sd x24, 192(sp)
sd x25, 200(sp)
sd x26, 208(sp)
```

```

sd x27, 216(sp)
sd x28, 224(sp)
sd x29, 232(sp)
sd x30, 240(sp)
sd x31, 248(sp)

# save sepc
csrr t0, sepc
sd t0, 0(sp)

csrr t0, scause
mv a0, t0
csrr t0, sepc
mv a1, t0
call trap_handler

ld t0, 0(sp)
csrw sepc, t0

ld x1, 8(sp)
ld x3, 24(sp)
ld x4, 32(sp)
ld x5, 40(sp)
ld x6, 48(sp)
ld x7, 56(sp)
ld x8, 64(sp)
ld x9, 72(sp)
ld x10, 80(sp)
ld x11, 88(sp)
ld x12, 96(sp)
ld x13, 104(sp)
ld x14, 112(sp)
ld x15, 120(sp)
ld x16, 128(sp)
ld x17, 136(sp)
ld x18, 144(sp)
ld x19, 152(sp)
ld x20, 160(sp)
ld x21, 168(sp)
ld x22, 176(sp)
ld x23, 184(sp)
ld x24, 192(sp)
ld x25, 200(sp)
ld x26, 208(sp)
ld x27, 216(sp)
ld x28, 224(sp)
ld x29, 232(sp)
ld x30, 240(sp)
ld x31, 248(sp)

addi sp, sp, context_size
sret

```

3.2.3 实现 trap 处理函数

在 `trap.c` 中实现 trap 处理函数 `trap_handler()`，其接收的两个参数分别是 `scause` 和 `sepc` 两个寄存器中的值。

```
void trap_handler(unsigned long scause, unsigned long sepc) {
    if ((long) scause < 0 && (scause & ((11 << 63) - 1)) == 5) {
        printk("%s", "Get STI!\n");
        clock_set_next_event();
    }
}
```

3.2.4 实现时钟中断相关函数

1. 在 `clock.c` 中实现 `get_cycles()`：使用 `rdtime` 汇编指令获得当前 `time` 寄存器中的值。
2. 在 `clock.c` 中实现 `clock_set_next_event()`：调用 `sbi_ecall`，设置下一个时钟中断事件。

```
unsigned long get_cycles() {
    unsigned long m_time;
    __asm__ volatile (
        "rdtime t0\n"
        "mv %[m_time], t0\n"
        : [m_time] "=r" (m_time)
    );
    return m_time;
}

void clock_set_next_event() {
    unsigned long next = get_cycles() + TIMECLOCK;

    sbi_ecall(0x0, 0x0, next, 0, 0, 0, 0, 0);
}
```

3.2.5 编译及测试

```
// test.c
void test() {
    for(int i = 1; i <= 120000000; ++i) {
        if(i == 120000000) {
            printk("%s", "kernel is running!\n");
            i = 0;
        }
    }
}
```



```

3 /* kernel代码起始位置 */
4 BASE_ADDR = 0x80200000;
5
6 查找命令
7
8 SECTIONS
9 {
10     /* . 代表当前地址 */
11     . = BASE_ADDR;
12
13     /* 记录kernel代码的起始地址 */
14     _kernel = .;
15     .set ignorecase 忽略大小写的查找
16
17     /* ALIGN(0x1000) 表示4KB对齐 */
18     /* _stext, _etext 分别记录了text段的起始与结束地址 */
19     .text : ALIGN(0x1000)
20     _stext = .;
21
22     /* 加入了 .text.init */
23     /* 之后我们实现 中断处理逻辑 会放置在 .text.entry */
24     .text.entry :
25     .set nohlsearch 关闭高亮搜索显示
26
27     /* 关闭当前的高亮显示, 如果再次输入 */
28     _etext = .;
29     .set hlinsearch 逐步搜索模式, 对当前输入的字
30
31     /* 重新搜索, 在搜索到文件头或尾 */
32     .rodata : ALIGN(0x1000)
33     _srodata = .;
34
35     替换命令
36     /*(.rodata .rodata.*)
37     将当前字替换为 当前字前加下划线
38
39     7 0000000008200318 t $x
40     8 000000000820035c t $x
41     9 00000000082003dc t $x
42     0 000000000820042c t $x
43     1 0000000008200908 t $x
44     2 0000000008200000 A BASE_ADDR
45     3 0000000008203000 B boot_stack
46     4 0000000008204000 B boot_stack_top
47     5 00000000082001b4 t clock_set_next_event
48     6 0000000008204000 B _ebss
49     7 0000000008202008 D _edata
50     8 0000000008204000 B _ekernel
51     9 00000000082010bc R _erodata
52     0 0000000008200908 T _etext单词词
53     1 0000000000000100 a frameize
54     2 0000000008200188 t get_cycles
55     3 0000000008202008 d _GLOBAL_OFFSET_TABLE_
56     4 0000000008200908 t printk
57     5 00000000082003dc t putc
58     6 0000000008200218 t sbi_ecall
59     7 0000000008203000 B _sbss
60     8 0000000008202000 D _sdata
61     9 0000000008200000 T _skernel
62     0 0000000008201000 R _srodata
63     1 0000000008200000 T _start
64     2 0000000008200318 t start_kernel
65     3 0000000008200000 t _stext
66     4 000000000820035c t test
67     5 0000000008202000 D TIMELOCK

```

3. 用 `csr_read` 宏读取 `sstatus` 寄存器的值，对照 RISC-V 手册解释其含义（截图）。

修改 test.c

```
#include <defs.h>
#include <printk.h>

struct sstatus {
    uint64 wpri_0 : 1;
    uint64 sie : 1;
    uint64 wpri_1 : 3;
    uint64 spie : 1;
    uint64 ube : 1;
    uint64 wpri_2 : 1;
    uint64 spp : 1;
    uint64 vs : 2;
    uint64 wpri_3 : 2;
    uint64 fs : 2;
    uint64 xs_l : 1;
    uint64 xs_h : 1;
    uint64 wpri_4 : 1;
    uint64 sum : 1;
    uint64 mxr : 1;
    uint64 wpri_5 : 4;
    uint64 wpri_6 : 8;
    uint64 uxl : 2;
    uint64 wpri_7 : 6;
    uint64 wpri_8 : 8;
    uint64 wpri_9 : 8;
    uint64 wpri_10 : 7;
    uint64 sd : 1;
}__attribute__((packed));

void test() {
    uint64 sstatus v = csr_read(sstatus);
```

```

struct sstatus *sstatus_0 = &sstatus_v;
printk("sstatus: \n");
printk("sie %lld\n", sstatus_0->sie);
printk("spie %lld\n", sstatus_0->spie);
printk("ube %lld\n", sstatus_0->ube);
printk("spp %lld\n", sstatus_0->spp);
printk("vs %lld\n", sstatus_0->vs);
printk("fs %lld\n", sstatus_0->fs);
printk("xs %lld\n", sstatus_0->xs_l | (sstatus_0->xs_h << 1));
printk("sum %lld\n", sstatus_0->sum);
printk("mxr %lld\n", sstatus_0->mxr);
printk("uxl %lld\n", sstatus_0->uxl);
printk("sd %lld\n", sstatus_0->sd);
csr_write(sscratch, 0x57);
printk("sscratch: 0x%x\n", csr_read(sscratch));
for(int i = 1; i <= 120000000; ++i) {
    if(i == 120000000) {
        printk("kernel is running!\n");
        i = 0;
    }
}
}

```

```

Boot HART MEDELEG          : 0x000000000000b109
2022 Hello RISC-V
sstatus:
sie 1
spie 0
ube 0
spp 0
vs 0
fs 3
xs 0
sum 0
mxr 0
uxl 0
sd 1
sscratch: 0x00000057
kernel is running!
Get STI!

```

如图 `csr_read` 和 `csr_write` 的实现正确。

【sie = 1】允许中断。

【spie = 0】进入S mode之前不允许中断。

【ube = 0】小端内存访问。

【spp = 0】之前的mode是U mode。

【vs = 0】禁止扩展S mode中虚拟内存功能。

【sum = 0】不允许用户访问内存。

【mxr = 0】不允许执行从用户模式内存读取的指令。

【sd = 1】允许S mode下的中断在S mode中处理。

4. 用 `csr_write` 宏向 `sscratch` 寄存器写入数据，并验证是否写入成功（截图）。

见上一题。

5. Detail your steps about how to get `arch/arm64/kernel/sys.i`

在 `/linux-6.6-rc2` 目录下执行：

```
sudo apt install g++-aarch64-linux-gnu binutils-aarch64-linux-gnu
make ARCH=arm64 defconfig
make arch/arm64/kernel/sys.i ARCH=arm64 CROSS_COMPILE=aarch64-linux-
gnu-
```

```
pac@ubuntu:~/Documents/linux-6.6-rc2$ make arch/arm64/kernel/sys.i ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu-
CALL scripts/checksyscalls.sh 将字符替换为a，当前字符即光标所在字符。
CPP arch/arm64/kernel/sys.i I
```

6. Find system call table of Linux v6.0 for `ARM32`, `RISC-V(32 bit)`, `RISC-V(64 bit)`, `x86(32 bit)`, `x86_64`

List source code file, the whole system call table with macro expanded, screenshot every step.

- `ARM32`: `arch/arm/kernel/entry-common.S`

```
syscall_table_start sys_call_table
#ifdef CONFIG_AEABI
#include <calls-eabi.S>
#else
#include <calls-oabi.S>
#endif
syscall_table_end sys_call_table
```

`arch/arm/include/generated/calls-eabi.S`

```
arch > arm > include > generated > asm calls-eabi.S
```

```
1  __SYSCALL(0, sys_restart_syscall)
2  __SYSCALL(1, sys_exit)
3  __SYSCALL(2, sys_fork)
4  __SYSCALL(3, sys_read)
5  __SYSCALL(4, sys_write)
6  __SYSCALL(5, sys_open)
7  __SYSCALL(6, sys_close)
8  __SYSCALL(7, sys_ni_syscall)
9  __SYSCALL(8, sys_creat)
10 __SYSCALL(9, sys_link)
11 __SYSCALL(10, sys_unlink)
12 __SYSCALL(11, sys_execve)
13 __SYSCALL(12, sys_chdir)
14 __SYSCALL(13, sys_ni_syscall)
15 __SYSCALL(14, sys_mknod)
16 __SYSCALL(15, sys_chmod)
17 __SYSCALL(16, sys_lchown16)
18 __SYSCALL(17, sys_ni_syscall)
19 __SYSCALL(18, sys_ni_syscall)
20 __SYSCALL(19, sys_lseek)
21 __SYSCALL(20, sys_getpid)
22 __SYSCALL(21, sys_mount)
23 __SYSCALL(22, sys_ni_syscall)
24 __SYSCALL(23, sys_setuid16)
25 __SYSCALL(24, sys_getuid16)
26 __SYSCALL(25, sys_ni_syscall)
27 __SYSCALL(26, sys_ptrace)
28 __SYSCALL(27, sys_ni_syscall)
29 __SYSCALL(28, sys_ni_syscall)
30 __SYSCALL(29, sys_pause)
31 __SYSCALL(30, sys_ni_syscall)
32 __SYSCALL(31, sys_ni_syscall)
33 __SYSCALL(32, sys_ni_syscall)
34 __SYSCALL(33, sys_access)
35 __SYSCALL(34, sys_nice)
```

由于 AS 没有预处理选项，所以无法进行宏展开。

- RISC-V(32 bit)

```
make ARCH=riscv 32-bit.config
make arch/riscv/kernel/syscall_table.i ARCH=riscv CROSS_COMPILE
=riscv64-linux-gnu-
```

```
arch/riscv/kernel/syscall_table.c
```

```
arch > riscv > kernel > C syscall_table.c
13 #define __SYSCALL(nr, call) asm linkage long __riscv_##call(const struct pt_regs *);
14 #include <asm/unistd.h>
15
16 #undef __SYSCALL
17 #define __SYSCALL(nr, call) [nr] = __riscv_##call,
18
19 void * const sys_call_table[__NR_syscalls] = {
20     [0 ... __NR_syscalls - 1] = __riscv_sys_ni_syscall,
21     #include <asm/unistd.h>
22 };
23
```

```
arch/riscv/kernel/syscall_table.i
```

```

arch > riscv > kernel > C syscall_table.i
58516 void * const sys_call_table[453] = {
58517     [0 ... 453 - 1] = __riscv_sys_ni_syscall,
58518     # 1 "./arch/riscv/include/asm/unistd.h" 1
58519     # 24 "./arch/riscv/include/asm/unistd.h"
58520     # 1 "./arch/riscv/include/uapi/asm/unistd.h" 1
58521     # 26 "./arch/riscv/include/uapi/asm/unistd.h"
58522     # 1 "./include/uapi/asm-generic/unistd.h" 1
58523     # 34 "./include/uapi/asm-generic/unistd.h"
58524     [0] = __riscv_sys_io_setup,
58525
58526     [1] = __riscv_sys_io_destroy,
58527
58528     [2] = __riscv_sys_io_submit,
58529
58530     [3] = __riscv_sys_io_cancel,
58531

```

- RISC-V(64 bit)

```

make ARCH=riscv 64-bit.config
make arch/riscv/kernel/syscall_table.i ARCH=riscv CROSS_COMPILE
=riscv64-linux-gnu-

```

arch/riscv/kernel/syscall_table.i

```

arch > riscv > kernel > C syscall_table.i
59064
59065
59066
59067
59068 void * const sys_call_table[453] = {
59069     [0 ... 453 - 1] = __riscv_sys_ni_syscall,
59070     # 1 "./arch/riscv/include/asm/unistd.h" 1
59071     # 24 "./arch/riscv/include/asm/unistd.h"
59072     # 1 "./arch/riscv/include/uapi/asm/unistd.h" 1
59073     # 26 "./arch/riscv/include/uapi/asm/unistd.h"
59074     # 1 "./include/uapi/asm-generic/unistd.h" 1
59075     # 34 "./include/uapi/asm-generic/unistd.h"
59076     [0] = __riscv_sys_io_setup,
59077
59078     [1] = __riscv_sys_io_destroy,
59079
59080     [2] = __riscv_sys_io_submit,
59081
59082     [3] = __riscv_sys_io_cancel,

```

- x86 (32 bit)

```

make ARCH=x86 i386_defconfig
make arch/x86/um/sys_call_table_32.i CROSS_COMPILE= ARCH=x86

```

arch/x86/um/sys_call_table_32.c

```
arch > x86 > um > C sys_call_table_64.c
24 #undef __SYSCALL
25 #define __SYSCALL(nr, sym) sym,
26
27 extern asmlinkage Long sys_ni_syscall(unsigned Long, unsigned Long, unsigned Long, unsigned Long, unsigned Long, unsigned Long);
28
29 const sys_call_ptr_t sys_call_table[] __cacheline_aligned = {
30 #include <asm/syscalls_64.h>
31 };
32
33 int syscall_table_size = sizeof(sys_call_table);
34
```

arch/x86/um/sys_call_table_32.i

```
arch > x86 > um > C sys_call_table_32.i
23838 extern __attribute__((regparm(0))) Long sys_ni_syscall(unsigned Long, unsigned Long, unsigned Long,
23839 Long);
23840 const sys_call_ptr_t sys_call_table[] __attribute__((__aligned__((1 << (5)))))) = {
23841 # 1 "./arch/x86/include/generated/asm/syscalls_32.h" 1
23842 sys_restart_syscall,
23843 sys_exit,
23844 sys_fork,
23845 sys_read,
23846 sys_write,
23847 sys_open,
23848 sys_close,
23849 sys_waitpid,
23850 sys_creat,
23851 sys_link,
```

- x86_64

```
make ARCH=x86 x86_64_defconfig
make arch/x86/um/sys_call_table_64.i CROSS_COMPILE= ARCH=x86
```

arch/x86/um/sys_call_table_64.c

```
arch > x86 > um > C sys_call_table_64.c
24 #undef __SYSCALL
25 #define __SYSCALL(nr, sym) sym,
26
27 extern asmlinkage Long sys_ni_syscall(unsigned Long, unsigned Long, unsigned Long, unsigned Long, unsigned Long, unsigned Long);
28
29 const sys_call_ptr_t sys_call_table[] __cacheline_aligned = {
30 #include <asm/syscalls_64.h>
31 };
32
33 int syscall_table_size = sizeof(sys_call_table);
34
```

arch/x86/um/sys_call_table_64.i

```
arch > x86 > um > C sys_call_table_64.i
23657
23658 const sys_call_ptr_t sys_call_table[] __attribute__((__aligned__((1 << (6)))))) = {
23659 # 1 "./arch/x86/include/generated/asm/syscalls_64.h" 1
23660 sys_read,
23661 sys_write,
23662 sys_open,
23663 sys_close,
23664 sys_newstat,
23665 sys_newfstat,
23666 sys_newlstat,
23667 sys_poll,
23668 sys_lseek,
23669 sys_mmap,
23670 sys_mprotect,
23671 sys_munmap,
```

7. Explain what is ELF file? Try readelf and objdump command on an ELF file, give screenshot of the output.
Run an ELF file and cat /proc/PID/maps to give its memory layout.

ELF 包含将序加载到内存中所必要的程序内存布局的数据结构（如程序头表、符号表、节头表）和各个段的具体数据。

如下为读取 **ELF** 头的截图：

```
litrehinn@litrehinn-soft-router:~/zzx/tmp$ readelf -h zzx
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                                   2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  DYN (Position-Independent Executable file)
  Machine:                                Advanced Micro Devices X86-64
  Version:                                0x1
  Entry point address:                    0x1060
  Start of program headers:               64 (bytes into file)
  Start of section headers:              13976 (bytes into file)
  Flags:                                  0x0
  Size of this header:                    64 (bytes)
  Size of program headers:                56 (bytes)
  Number of program headers:              13
  Size of section headers:                64 (bytes)
  Number of section headers:              31
  Section header string table index:      30
```

如下为读取 **ELF** 符号表的截图：

```
litrehinn@litrehinn-soft-router:~/zzx/tmp$ readelf -s zzx
Symbol table '.dynsym' contains 7 entries:
Num:  Value              Size Type Bind Vis Ndx Name
 0: 0000000000000000      0 NOTYPE LOCAL DEFAULT UND
 1: 0000000000000000      0 FUNC GLOBAL DEFAULT UND [...]@GLIBC_2.34 (2)
 2: 0000000000000000      0 NOTYPE WEAK DEFAULT UND _ITM_deregisterT[...]
 3: 0000000000000000      0 FUNC GLOBAL DEFAULT UND puts@GLIBC_2.2.5 (3)
 4: 0000000000000000      0 NOTYPE WEAK DEFAULT UND __gmon_start__
 5: 0000000000000000      0 NOTYPE WEAK DEFAULT UND _ITM_registerTMC[...]
 6: 0000000000000000      0 FUNC WEAK DEFAULT UND [...]@GLIBC_2.2.5 (3)

Symbol table '.symtab' contains 36 entries:
Num:  Value              Size Type Bind Vis Ndx Name
 0: 0000000000000000      0 NOTYPE LOCAL DEFAULT UND
 1: 0000000000000000      0 FILE LOCAL DEFAULT ABS Scrt1.o
 2: 0000000000000038c      32 OBJECT LOCAL DEFAULT 4 __abi_tag
 3: 0000000000000000      0 FILE LOCAL DEFAULT ABS crtstuff.c
 4: 0000000000001090      0 FUNC LOCAL DEFAULT 16 deregister_tm_clones
 5: 00000000000010c0      0 FUNC LOCAL DEFAULT 16 register_tm_clones
 6: 0000000000001100      0 FUNC LOCAL DEFAULT 16 __do_global_ctors_aux
 7: 0000000000003d00      1 OBJECT LOCAL DEFAULT 26 completed.0
 8: 0000000000003dc0      0 OBJECT LOCAL DEFAULT 22 __do_global_dtor[...]
 9: 0000000000001140      0 FUNC LOCAL DEFAULT 16 frame_dummy
10: 0000000000003db8      0 OBJECT LOCAL DEFAULT 21 __frame_dummy_in[...]
11: 0000000000000000      0 FILE LOCAL DEFAULT ABS zzx.c
12: 0000000000000000      0 FILE LOCAL DEFAULT ABS crtstuff.c
13: 00000000000020f0      0 OBJECT LOCAL DEFAULT 20 __FRAME_END__
14: 0000000000000000      0 FILE LOCAL DEFAULT ABS
15: 0000000000003dc8      0 OBJECT LOCAL DEFAULT 23 _DYNAMIC
16: 0000000000002010      0 NOTYPE LOCAL DEFAULT 19 __GNU_EH_FRAME_HDR
17: 0000000000003fb8      0 OBJECT LOCAL DEFAULT 24 _GLOBAL_OFFSET_TABLE_
18: 0000000000000000      0 FUNC GLOBAL DEFAULT UND __libc_start_mai[...]
19: 0000000000000000      0 NOTYPE WEAK DEFAULT UND _ITM_deregisterT[...]
20: 0000000000004000      0 NOTYPE WEAK DEFAULT 25 data_start
21: 0000000000000000      0 FUNC GLOBAL DEFAULT UND puts@GLIBC_2.2.5
22: 0000000000004010      0 NOTYPE GLOBAL DEFAULT 25 _edata
23: 0000000000001168      0 FUNC GLOBAL HIDDEN 17 _fini
24: 0000000000004000      0 NOTYPE GLOBAL DEFAULT 25 __data_start
25: 0000000000000000      0 NOTYPE WEAK DEFAULT UND __gmon_start__
26: 0000000000004008      0 OBJECT GLOBAL HIDDEN 25 __dso_handle
27: 0000000000002000      4 OBJECT GLOBAL DEFAULT 18 _IO_stdin_used
28: 0000000000004018      0 NOTYPE GLOBAL DEFAULT 26 _end
29: 0000000000001060      38 FUNC GLOBAL DEFAULT 16 _start
30: 0000000000004010      0 NOTYPE GLOBAL DEFAULT 26 __bss_start
31: 0000000000001149      30 FUNC GLOBAL DEFAULT 16 main
32: 0000000000004010      0 OBJECT GLOBAL HIDDEN 25 __TMC_END__
33: 0000000000000000      0 NOTYPE WEAK DEFAULT UND _ITM_registerTMC[...]
34: 0000000000000000      0 FUNC WEAK DEFAULT UND __cxa_finalize@G[...]
35: 0000000000001000      0 FUNC GLOBAL HIDDEN 12 _init
```

如下为读取 **ELF** 程序头表（各段的信息）的截图：


```

litrehin@litrehin-soft-router:~/zxx/tmp$ readelf -l zxx
Elf file type is DYN (Position-Independent Executable file)
Entry point 0x1060
There are 13 program headers, starting at offset 64

Program Headers:
Type           Offset             VirtAddr           PhysAddr
               FileSiz            MemSiz             Flags   Align
PHDR           0x0000000000000000 0x0000000000000000 0x0000000000000000
INTERP         0x0000000000000118 0x0000000000000118 0x0000000000000118
               0x000000000000001c 0x000000000000001c  R      0x1
               [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
LOAD           0x0000000000000000 0x0000000000000000 0x0000000000000000
               0x0000000000000028 0x0000000000000028  R      0x1000
LOAD           0x0000000000000100 0x0000000000000100 0x0000000000000100
               0x0000000000000175 0x0000000000000175  R E    0x1000
LOAD           0x0000000000000200 0x0000000000000200 0x0000000000000200
               0x00000000000000f4 0x00000000000000f4  R      0x1000
LOAD           0x00000000000002b8 0x00000000000003b8 0x00000000000003b8
               0x0000000000000258 0x0000000000000260  RW     0x1000
DYNAMIC         0x00000000000002c8 0x00000000000003c8 0x00000000000003c8
               0x00000000000001f0 0x00000000000001f0  RW     0x8
NOTE           0x0000000000000338 0x0000000000000338 0x0000000000000338
NOTE           0x0000000000000360 0x0000000000000368 0x0000000000000368
               0x00000000000000u4 0x00000000000000u4  R      0x4
GNU_PROPERTY   0x0000000000000338 0x0000000000000338 0x0000000000000338
GNU_EH_FRAME   0x0000000000000330 0x0000000000000330  R      0x8
GNU_STACK      0x0000000000000000 0x0000000000000000 0x0000000000000000
GNU_RELRO      0x00000000000002b8 0x00000000000003b8 0x00000000000003b8
               0x00000000000002b8 0x00000000000002b8  R      0x1

Section to Segment mapping:
Segment Sections...
00
01  .interp
02  .interp.note.gnu.property.note.gnu.build-id.note.ABI-tag.gnu.hash.dynsym.dynstr.gnu.version.gnu.version_r.rela.dyn.rela.plt
03  .init.plt.plt.got.plt.sec.text.fini
04  .rodata.eh_frame_hdr.eh_frame
05  .init_array.fini_array.dynamic.got.data.bss
06  .dynamic
07  .note.gnu.property
08  .note.gnu.build-id.note.ABI-tag
09  .note.gnu.property
10  .eh_frame_hdr
11
12  .init_array.fini_array.dynamic.got

```

如下为读取 **ELF** 各节的截图：

```

litrehin@litrehin-soft-router:~/zxx/tmp$ readelf -S zxx.o
There are 14 section headers, starting at offset 0x258:

Section Headers:
[Nr] Name                               Type                               Address                               Offset
     Size                               EntSize                            Flags Link Info Align
[ 0]                                     NULL                               0000000000000000 00000000
     0000000000000000 0000000000000000 0 0 0
[ 1] .text                                PROGBITS                           0000000000000000 00000040
     000000000000001e 0000000000000000 AX 0 0 1
[ 2] .rela.text                           RELA                               0000000000000000 00000198
     0000000000000030 0000000000000018 I 11 1 8
[ 3] .data                                PROGBITS                           0000000000000000 0000005e
     0000000000000000 0000000000000000 WA 0 0 1
[ 4] .bss                                NOBITS                             0000000000000000 0000005e
     0000000000000000 0000000000000000 WA 0 0 1
[ 5] .rodata                              PROGBITS                           0000000000000000 0000005e
     000000000000000c 0000000000000000 A 0 0 1
[ 6] .comment                             PROGBITS                           0000000000000000 0000006a
     000000000000002c 0000000000000001 MS 0 0 1
[ 7] .note.GNU-stack                       PROGBITS                           0000000000000000 00000096
     0000000000000000 0000000000000000 0 0 1
[ 8] .note.gnu.pr[...]                     NOTE                               0000000000000000 00000098
     0000000000000020 0000000000000000 A 0 0 8
[ 9] .eh_frame                             PROGBITS                           0000000000000000 000000b8
     0000000000000038 0000000000000000 A 0 0 8
[10] .rela.eh_frame                         RELA                               0000000000000000 000001c8
     0000000000000018 0000000000000018 I 11 9 8
[11] .symtab                               SYMTAB                             0000000000000000 000000f0
     0000000000000090 0000000000000018 12 4 8
[12] .strtab                               STRTAB                             0000000000000000 00000180
     0000000000000011 0000000000000000 0 0 1
[13] .shstrtab                             STRTAB                             0000000000000000 000001e0
     0000000000000074 0000000000000000 0 0 1

Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
D (mbind), l (large), p (processor specific)

```

对 **zxx.o** 进行反汇编得到的汇编代码：

对 `zzx.o` 的各节进行解析得到的二进制 dump:

执行 **ELF** 文件，并输出其内存布局：

8. 通过查看 RISC-V Privileged Spec 中的 `medeleg` 和 `mideleg`，解释上面 `MIDELEG` 值的含义。

No. 18 / 19

Interrupt	Exception Code	Description
1	0	<i>Reserved</i>
1	1	Supervisor software interrupt
1	2	<i>Reserved</i>
1	3	Machine software interrupt
1	4	<i>Reserved</i>
1	5	Supervisor timer interrupt
1	6	<i>Reserved</i>
1	7	Machine timer interrupt
1	8	<i>Reserved</i>
1	9	Supervisor external interrupt

MIDELEG 值 0222:

- MIDELEG[1] = 1 表示把 Supervisor software interrupt 委托给 S-mode
- MIDELEG[5] = 1 表示把 Supervisor timer interrupt 委托给 S-mode
- MIDELEG[9] = 1 表示把 Supervisor external interrupt 委托给 S-mode